

Open Tiled Manycore System-on-Chip*

Stefan Wallentowitz Philipp Wagner
Michael Tempelmeier Thomas Wild
Andreas Herkersdorf

Institute for Integrated Systems, TU München
Arcisstr. 21, D-80290 München
`stefan.wallentowitz@tum.de`

April 19, 2013

Abstract

Manycore System-on-Chip include an increasing amount of processing elements and have become an important research topic for improvements of both hardware and software. While research can be conducted using system simulators, prototyping requires a variety of components and is very time consuming. With the *Open Tiled Manycore System-on-Chip (OpTiMSoC)* we aim at building such an environment for use in our and other research projects as prototyping platform.

This paper describes the project goals and aspects of OpTiMSoC and summarizes the current status and ideas.

1 Motivation & Project Goals

Multicore System-on-Chip have become dominant in research and industry as high-performance and power-efficient processing platforms. With increasing amounts of processing elements the tile organization is a popular way to organize the elements of such a platform. Those platforms base on Network-on-Chip with a regular organization, often a mesh. Processing elements, memories and I/O devices are connected to this interconnect as *tiles*. This allows for platform creation by replicating these basic tiles to larger platforms. Some examples for tiled platforms are Tiler's processors or Intel's "Single Chip Cloud Computer".

Many aspects of such platforms are currently in the focus of research, for example system design aspects, communication-to-computation coupling, coherency and consistency issues, programming of future massively parallel platforms and many more.

Research of improvements of future manycore system-on-chip is mainly performed using system simulation approaches, where a variety of tools exist. Nevertheless when it comes to prototyping, hardware architects rely on a platform

*Dagstuhl Seminar 13052, "Multicore Enablement for Embedded and Cyber Physical Systems", organizers: Andreas Herkersdorf, Michael G. Hinchey, Michael Paulitsch, 27.01.13–01.02.13

they can prototype their ideas with. Building such a platform from scratch is very time consuming and a complete – not even speaking of common – platform has not been established. OpTiMSoC tries to fill this gap by providing the necessary environment to augment research prototyping in the respective field.

The goals of OpTiMSoC can therefore be summarized as follows:

- Build a foundation library of hardware elements that can shape a variety of tiled manycore platforms
- Support several different target platforms, ranging from simulation to FPGA-based emulation
- Provide the required infrastructure to compose and build those platforms
- Include a programming environment and runtime system
- Enable debugging with trace-based debugging techniques with a common middleware layer

Along those goals this paper will present the current status of the OpTiMSoC project. The current development of the project can be found on the project website¹. The project is open to contributions and aims to always provide paths that do not require licenses for necessary tools.

2 Basic Hardware Elements

The central element of OpTiMSoC is LISNoC², a basic **Network-on-Chip** implementation we developed. The basic NoC is a packet-switched, wormhole-forwarding, buffered implementation supporting virtual channels to avoid message dependent deadlocks. Variations supporting priorities, multicasts or bufferless forwarding are also part of LISNoC.

The Network-on-Chip is the underlying structure for the whole platform. Mesh topologies are commonly deployed in tiled manycore System-on-Chip, but also other topologies such as rings or hierarchical structures can be easily implemented using the LISNoC hardware elements library.

Beside the communication foundation, the **processing elements** are the important elements in OpTiMSoC. The basic processing element used in OpTiMSoC is the OpenRISC processor³. Despite the current infrastructure is built around this processor core, other alternatives will be added in the future, for example the LEON3 Sparc implementation⁴.

Some peripheral elements for the tiles, such as memories, interconnect etc. are also part of the library, together with additional I/O elements. Those strongly depend on the target platform as described in the following.

Crucial elements, which are under research currently, are those bridging communication and computation, namely the **network adapter** or also network interface (NA/NI). The central elements of such network adapters are:

¹<http://www.optimsoc.org>

²<http://www.lisnoc.org>

³<http://www.opencores.org/openrisc>

⁴<http://www.gaisler.com/>

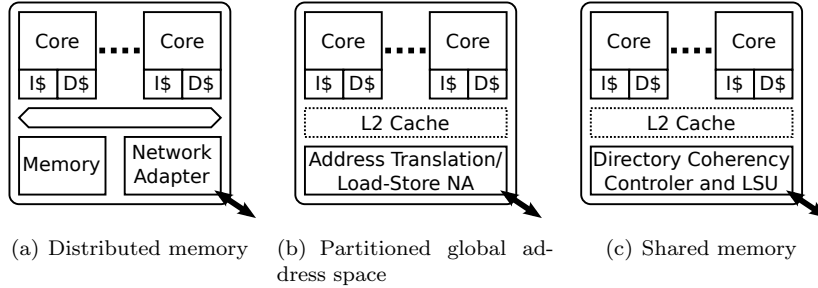


Figure 1: Different tile implementation alternatives

- Handle *memory transfers* between tiles and the memories.
- Provide hardware means to send *messages* between the tiles such as for processor communication or external data streams.

The services provided by a network adapter implementation depend on the system and **tile organization**. In OpTiMSoC we aim to cover three central organization styles as sketched in Figure 1. The choice of the organization strongly correlates with the choice of the programming model. At the moment we focus on distributed memory organization and message passing programming.

In distributed memory (see Figure 1(a)) a varying amount of processor cores is connected to a locally shared memory. Due to restricted memory sizes, a partitioned global address space variant (see Figure 1(b)) is used in prototyping platforms, where the network adapter employs an MPU-like load-store unit (LSU) that translates memory access addresses in a way that the global memory is partitioned in separate chunks, each assigned to one tile.

Finally, we currently work towards a platform variant which provides globally shared memory among all tiles (see Figure 1(c)). In the other platforms the locally shared memory or address space needs to be coherent, which is done with a write-through snooping policy. In future systems it is planned to augment this with a level 2 write-back directory-based cache coherency.

3 Target Platforms

With the current releases we try to cover a basic number of target platforms for development and prototyping of manycore System-on-Chip. In the following we describe the basic ways to implement an OpTiMSoC instance currently as they cover the main use cases.

RTL simulation using EDA tools is the basic entry point from a hardware designers point of view. For development of hardware elements in OpTiMSoC this step is essential. Nevertheless, most tools are commercial, so that a second variant (Verilator) is preferred in other cases.

Verilated simulation is a variant of RTL simulation. Verilator⁵ is a toolchain

⁵<http://www.veripool.org/wiki/verilator>

that compiles RTL code to C++ code for simulation. A SystemC⁶ module can easily be generated by Verilator. The main advantage of the use of Verilator over RTL simulation with (potentially costly) EDA tools is that it is completely open source. That allows software developers to develop code without prototyping hardware or the necessity for a commercial simulation tool.

Xilinx University Program Board 5 (XUPV5)⁷ is a widely used FPGA board often found in academic institutions. It employs a mid-range Virtex 5 FPGA and multiple I/O. Currently the DDR memory, UART and LCD display are used in OpTiMSoC.

ZTEX 1.15 boards⁸ are small outline boards including a Spartan 6 FPGA, DDR3 memory and a Cypress EZ-USB interface chip. The boards are easily used in standalone mode only, using a USB cable. Two small variants do not require a costly synthesis license but can be used with the XILINX WebPack. The boards are especially powerful using our debug infrastructure as described below and can be very useful to software developers.

CHIPit emulation platform⁹ is an emulation platform based on multiple FPGAs and allows for the prototyping of much larger systems. In the future OpTiMSoC will include a setup for such large systems on the emulator platform and the related infrastructure.

The target platforms described above are of course only a starting point. We aim to support many more targets in the future. Once the basic operation and I/O of other targets is supported (or is already worked out from other uses) it is relatively simple to include them in OpTiMSoC.

4 Platform Generator Tool

As introduced before, OpTiMSoC is not a platform proposal for one specific setup regarding dimensions or organization of the system. Instead it is intended as a basic set of elements and the necessary infrastructure. To allow the user to generate many different variants and configuration for a variety of targets easily, we previously presented our envisioned tool flow in [1].

The elements, tile organization and target platforms described above can be seen as the library elements to a platform generator tool. Based on these elements two separate platform generation steps are envisioned:

Platform description mapping: The platform is described based on generic layout patterns. It defines the interconnect structure, the organization of tiles in the system and the tile-internal organization. This description is augmented with some basic parameters, such as the processor implementation or similar. The output is an OpTiMSoC configuration.

⁶<http://www.accelera.org>

⁷<http://www.xilinx.com/univ/xupv5-1x110t.htm>

⁸<http://www.ztex.de>

⁹<http://www.synopsys.com>

OpTiMSoC configuration mapping: An OpTiMSoC configuration can be either described manually by the designer or can be the output of the platform description mapping. This mapping generates the actual files for the used elements for a given target and the respective build files, such as scripts or makefiles.

For a more detailed discussion we refer the reader to [1].

5 Programming and Runtime System

Previously we presented the hardware elements and targets of OpTiMSoC. Of same importance of course is the software part of OpTiMSoC, what involves the question how to program it and what runtime system it employs. Basically, two **underlying systems** are the fundamental base of OpTiMSoC software:

Baremetal system The baremetal system involves all drivers necessary to execute on the processing hardware, which involves the stack and heap management, exception handling etc. Therefore we essentially provide a port of the newlib libc implementation for OpTiMSoC based on the OpenRISC development (that also includes the gcc compiler).

Lean runtime system Based on the runtime system a simple runtime system is implemented that provides the central functions needed for more sophisticated systems: a thread scheduler and virtual memory management. Based on these, a simple microkernel or even other operating systems can be ported. It has to be noted that the runtime system, as most other elements of OpTiMSoC, is not optimized in a sense that we tried to get every cycle out of there, but was instead designed with the goals of readability and simpleness.

Both systems are augmented with the basic drivers for the hardware elements. Furthermore high-level **programming APIs** are required for parallel programming. We decided to provide baseline implementations of the Multicore Association APIs¹⁰ as they perfectly cover the problems of embedded systems. Therefore we currently work on the implementation of:

MCAPI The communication API is a basic message passing library used for inter-process communication. The API handles the communication over the Network-on-Chip abstracting from the employed network adapter capabilities. Different transport layer implementations instead handle the different hardware implementations.

MTAPI Recently the task management API was released. We work towards a distributed implementation of the MTAPI in OpTiMSoC. A special focus for our future work will be the handling of heterogenous platforms which include dedicated hardware accelerator tiles that offload the regular processing elements with compute intensive tasks, such as crypto or signal processing accelerators.

¹⁰<http://www.mutlicore-association.org>

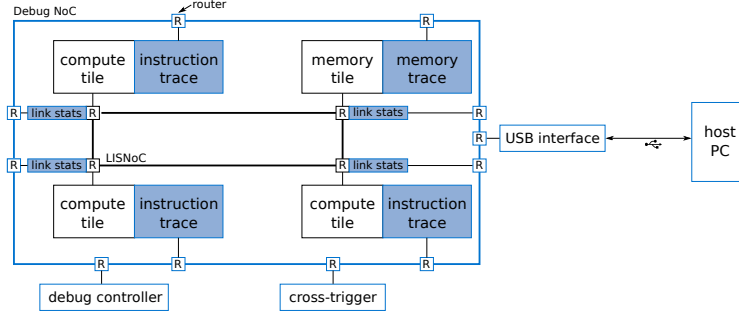


Figure 2: A sample 2x2 system with the debug infrastructure (blue).

6 Debugging and Diagnosis Infrastructure

Optimizing the system and writing software for OpTiMSoC, like for any other tightly integrated System-on-Chip, is complicated by the limited system observability. To overcome this problem a debugging and diagnosis infrastructure was integrated into OpTiMSoC. The most significant challenge of traditional run-control based debugging approaches, the scaling to heterogenous multi-core platforms with different clock domains, is resolved by a modular and decentralized tracing-based solution, as depicted in Figure 2 for an exemplary 2x2 system.

Each system component that should be observed can be extended with a specific debug module, which collects relevant data from it in the form of traces. The content of these traces and the definition of “relevant data” can differ greatly between the data sources; from CPU cores instruction traces representing the program flow are collected, from routers in the NoC aggregated link usage statistics, and from the memory data traces can be extracted (not yet part of OpTiMSoC).

Triggers can be set to reduce the amount of data collected. A trigger is a condition that can cause the collection of trace data to start or to stop. These conditions can also be combined across different debug modules (cross-triggers).

All collected data is then augmented with a timestamp to enable a temporal correlation of trace messages from different sources, possibly compressed¹¹ and then sent over an independent, 16-bit wide buffered Ring-NoC (the so-called Debug NoC) to an external interface connected to a host PC.

Depending on the target hardware different interfaces can be used to transfer the tracing data off-chip. Currently implemented are USB 2.0 for the ZTEX 1.15 boards and TCP for OpTiMSoC running as ModelSim simulation.

On the host PC side, a generic library, liboptimsosochost, abstracts from the different transport layers and provides a high-level interface to applications. Current user applications are optimsoc_cli and a basic graphical application.

¹¹Currently compression is only implemented for instruction traces.

7 Project Status & Roadmap

A young and large project like OpTiMSoC constantly evolves in many directions. At this time we work on integrating and documenting the aspects covered in this paper plus some special topics, such as support for different clock domains with the appropriate clock-domain crossings and clock management, the debug system and other components that have been already prepared in-house. The releases of different elements is spread over the second quarter of 2013.

The focus of future work is directed towards the challenges of a shared memory implementation (cache organization and coherence), the integration of more hardware accelerator options, and novel approaches in the field of debugging and diagnosis as well as in the runtime support system.

Apart from our own roadmap we hope to be able to integrate contributions from others which might evolve OpTiMSoC in yet unknown directions. We are glad to share our work and hope it helps other researchers in their work and welcome any feedback or contributions.

References

- [1] S. Wallentowitz, A. Lankes, A. Zaib, T. Wild, and A. Herkersdorf. A framework for open tiled manycore system-on-chip. In *Field Programmable Logic and Applications (FPL), 2012 22nd International Conference on*, pages 535–538, 2012. <http://dx.doi.org/10.1109/FPL.2012.6339273>.